


# Random Forest Challenge

December 2, 2020

```
In [1]: import numpy as np # to build the algorithm
import matplotlib.pyplot as plt # to visualize
from sklearn.datasets import make_circles # to generate a dataset

In [2]: # Generate a dataset
X, y = make_circles(n_samples=100, noise=0.1, factor=0.5, random_state=0)
plt.figure(dpi=200)
plt.scatter(X[:, 0][y == 0], X[:, 1][y == 0], label=0)
plt.scatter(X[:, 0][y == 1], X[:, 1][y == 1], label=1)
plt.xlabel('feature 0')
plt.ylabel('feature 1')
plt.legend()
```

Out[2]: <matplotlib.legend.Legend at 0x11ca3def0>



```

In [3]: # Train the model
        from sklearn.ensemble import RandomForestClassifier

        clf = RandomForestClassifier(n_estimators=30)
        clf.fit(X, y)

Out[3]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                               max_depth=None, max_features='auto', max_leaf_nodes=None,
                               min_impurity_decrease=0.0, min_impurity_split=None,
                               min_samples_leaf=1, min_samples_split=2,
                               min_weight_fraction_leaf=0.0, n_estimators=30, n_jobs=None,
                               oob_score=False, random_state=None, verbose=0,
                               warm_start=False)

In [4]: # Test the model
        X_test, y_test = make_circles(n_samples=100, noise=0.1, factor=0.5, random_state=42)
        y_pred = clf.predict(X_test)
        acc = sum(y_pred == y_test)/len(y_test)
        print('Accuracy: ', acc)

Accuracy:  0.97

In [5]: clf.feature_importances_

Out[5]: array([0.48581498, 0.51418502])

In [6]: # Challenge 1
        def gini_calculator(y):
            """
            Calculates the gini impurity of a set
            Arguments
                y: np.array() containing the labels
            Returns
                gini:  $1 - p_0^2 - p_1^2$ 
                    p0 ratio of class 0
                    p1 ratio of class 1
            """

            return gini

In [7]: # Sanity check
        num_datapoints = 20
        plt.figure(dpi=200)
        for i in range(num_datapoints+1):
            num_ones = i


```

```

num_zeros = num_datapoints-i
p1 = num_ones/num_datapoints
combined_set = np.concatenate((np.ones(num_ones), np.zeros(num_zeros)))
gini = gini_calculator(combined_set)
plt.scatter(p1, gini, color='k')
plt.xlabel('Ratio of 1')
plt.ylabel('Gini impurity')

```

Out[7]: Text(0, 0.5, 'Gini impurity')



```

In [8]: def gini_of_a_split(y1, y2):
        '''
        Weighted average gini of two sets
        Arguments
            y1: np.array() containing the labels of set 1
            y2: np.array() containing the labels of set 2
        Returns
            avg_gini: Weighted average gini of y1 and y2
        '''
        g1 = gini_calculator(y1)
        w1 = len(y1)/(len(y1)+len(y2))
        g2 = gini_calculator(y2)
        w2 = len(y2)/(len(y1)+len(y2))
        avg_gini = g1*w1 + g2*w2
        return avg_gini

```

```

In [9]: # Challenge 2
        from operator import itemgetter

def split_finder(X, y):
    """
    Finds the best split that minimizes the average gini.
    Best split is defined by a feature index and its value.

    Arguments
        m = num_of_datapoints
        n = num_of_features
        X: np.array() shape (m, n)
        y: np.array() shape (m, 1)

    Returns
        best_split_feature: integer, best feature index
        best_split_value: float, best value

    """


    return best_split_feature, best_split_value

In [10]: # Visualize the first split
best_split = split_finder(X, y)
plt.figure(dpi=200)
plt.scatter(X[:, 0][y == 0], X[:, 1][y == 0], label=0, s=5)
plt.scatter(X[:, 0][y == 1], X[:, 1][y == 1], label=1, s=5)
plt.xlabel('feature 0')
plt.ylabel('feature 1')
plt.legend()

split_feature = best_split[0]
split_value = best_split[1]
boundary_limits = [min(X[:, 1]), max(X[:, 1])]

if split_feature == 0:
    plt.plot([split_value, split_value], boundary_limits, lw=.5)
if split_feature == 1:
    plt.plot(boundary_limits, [split_value, split_value], lw=.5)

```



```
In [11]: def splitter(X, y):
```

```
    '''
```

*This is one node split i.e. building block of a tree.*

*Given X and y,*

- finds the split*
- splits the datasets into 2 subsets*
- returns the subsets and the split*

*Arguments*

```
    m = num_of_datapoints
    n = num_of_features
    X: np.array() shape (m, n)
    y: np.array() shape (m, 1)
```

*Returns*

```
    subset1: list of np arrays, [X1, y1]
             X1: np.array() shape (m1, n)
             y1: np.array() shape (m1, 1)
    subset2: list of np arrays, [X2, y2]
             X2: np.array() shape (m2, n)
             y2: np.array() shape (m2, 1)
    where m = m1 + m2
```

```
    split: a tuple, (split_feature, split_value)
```

```

        split_feature: integer, best feature index
        split_value: float, best value
    """
    split = split_finder(X, y)

    X1 = X[X[:, split[0]] < split[1]]
    y1 = y[X[:, split[0]] < split[1]]

    X2 = X[X[:, split[0]] > split[1]]
    y2 = y[X[:, split[0]] > split[1]]

    return (X1, y1), (X2, y2), split


In [12]: # Visualize the split and the subsets
subset1, subset2, split = splitter(X, y)
X1 = subset1[0]
X2 = subset2[0]

plt.figure(dpi=200)
plt.scatter(X1[:, 0], X1[:, 1], color='k', label='subset 1')
plt.scatter(X2[:, 0], X2[:, 1], color='c', label='subset 2')
plt.legend()

split_feature = split[0]
split_value = split[1]
boundary_limits = [min(X[:, 1]), max(X[:, 1])]

if split_feature == 0:
    plt.plot([split_value, split_value], boundary_limits, lw=.5)
if split_feature == 1:
    plt.plot(boundary_limits, [split_value, split_value], lw=.5)

```



In [13]: # Challenge 3

```
def fit_tree(X, y):
```

```
    '''
```

*Repeat the splitter to fit a tree to X and y.  
return the tree i.e. the trained model*

*Arguments*

*m = num\_of\_datapoints*

*n = num\_of\_features*

*X: np.array() shape (m, n)*

*y: np.array() shape (m, 1)*

*Returns*

*tree: a dictionary of nodes*

*key: node name e.g. 'root', 'rootRL'*

*value: node dictionary*

*key:*

*'depth' int*

*'data' (subset\_X, subset\_y)*

*'class' majority class*

*'leaf' binary, leaf(1) or not(0)*

*'split\_feature' 0 or 1*

*'split\_value' float*

```
    '''
```

```

        return tree

In [14]: tree = fit_tree(X, y)

In [15]: # Challenge 4
def predict_tree(X, tree):
    """
    Given X and the model (i.e. tree), return predictions

    Arguments
    X: np.array() shape (m, n)
    tree: a dictionary of nodes

    Returns
    y_pred: a list of predictions for each row of X,
            class labels 0 or 1.

    """

    return y_pred

In [16]: y_pred = predict_tree(X, tree)

In [17]: def accuracy(y_pred, y):
    return sum(y_pred == y)/len(y)

In [18]: accuracy(y_pred, y)

Out[18]: 1.0

In [19]: # Putting all together
# with Train/Test
X_train, y_train = make_circles(n_samples=100, noise=0.1, factor=0.5, random_state=0)
X_test, y_test = make_circles(n_samples=100, noise=0.1, factor=0.5, random_state=42)
tree = fit_tree(X_train, y_train)
y_pred_train = predict_tree(X_train, tree)
y_pred_test = predict_tree(X_test, tree)

print('Training acc:', accuracy(y_pred_train, y_train))
print('Testing acc:', accuracy(y_pred_test, y_test))

Training acc: 1.0
Testing acc: 0.94

In [24]: # Challenge 5
def fit_forest(X, y):

```



```

'''
Fit 30 trees
by randomly sampling from
X and y
return 30 trees

Arguments
    X: np.array() shape (m, n)
    y: np.array() shape (m, 1)

Returns
    forest: a list of trees

'''
num_trees = 30
forest = []
for i in range(num_trees):
    X_sample =
    y_sample =
    tree = fit_tree(X_sample, y_sample)
    forest.append(tree)
return forest

```

In [21]: # Challenge 6

```

def predict_forest(X, forest):
'''
Predict the labels for X
for all 30 trees
calculate the average of 30 trees
return avg. predictions

Arguments
    X: np.array() shape (m, n)
    forest: a list of trees

Returns
    y_pred: a list containing predicted classes
            for each row in X

'''

return y_pred

```

In [22]: # Putting all together

```

# with Train/Test
X_train, y_train = make_circles(n_samples=100, noise=0.1, factor=0.5, random_state=0)
X_test, y_test = make_circles(n_samples=100, noise=0.1, factor=0.5, random_state=42)
forest = fit_forest(X_train, y_train)

```

```
y_pred_train = predict_forest(X_train, forest)
y_pred_test = predict_forest(X_test, forest)

print('Training acc:', accuracy(y_pred_train, y_train))
print('Testing acc:', accuracy(y_pred_test, y_test))
```

```
Training acc: 1.0
Testing acc: 0.96
```

```
In [ ]:
```